

FIFA World Cup 2026 as a Contamination-Free Benchmark for LLM Forecasting Agents: Four Models, a Bookmaker, and 104 Matches

Jiacheng Ding*
University of Memphis
Memphis, Tennessee, USA
jding2@memphis.edu

Cong Guo*
University of Memphis
Memphis, Tennessee, USA
cguo@memphis.edu

Abstract

We introduce **WC2026-Agents**, a benchmark and dataset for evaluating large language models (LLMs) as *autonomous forecasting agents* on real, future events. For every one of the 104 matches of the 2026 FIFA World Cup, four frontier models — Claude Opus 4.8, ChatGPT (GPT-5.5, high reasoning), Gemini 3.1 Pro, and Grok (Expert Mode) — ran an identical search-act-reflect loop: gather evidence with a web tool, commit to a 1X2 (team-A win / draw / team-B win) probability distribution and a virtual \$100 bet, and, after the match, reflect given only the final score. Because every match kicked off after the models’ training cutoffs, the benchmark is *contamination-free* by construction. Crucially, we pair the four agents with a fifth competitor drawn from the same information environment — the pre-match betting market — collected as per-match 1X2 odds, giving an economically grounded human baseline and letting us score not just what an agent predicts but what it *does with money*. The release contains 416 forecasts and 414 reflections with verbatim reasoning, ground truth (including penalty shootouts), odds, and a reproducible evaluation suite. A reference evaluation surfaces findings that raw accuracy hides: the four agents issue an identical top pick in 92% of matches and none beats the market’s Brier score, yet they diverge sharply as decision-makers — betting return-on-investment ranges from -18% to +10%, fading the market is unprofitable for all four, the share of forecasts that cite the market ranges from 12% to 100%, and self-reported error rates on wrong picks range from 36% to 86%. The benchmark thus measures calibration, decision quality, and self-knowledge — axes on which frontier models differ even when their predictions do not. Data and code: <https://github.com/graphuofm/FIFA2026LLM>.

CCS Concepts

• Computing methodologies → Artificial intelligence; • Information systems → Data mining.

*Both authors contributed equally to this research.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
KDD ’27, San Jose, CA, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM
<https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

Keywords

large language models, LLM agents, forecasting, calibration, prediction markets, sports analytics, benchmark dataset, data contamination

1 Introduction

Large language models are moving from answering questions to *acting*: retrieving evidence, reasoning, committing to decisions, and reflecting on outcomes [22, 34, 36]. Evaluating this agentic behaviour is hard for three reasons. First, most benchmarks are static and increasingly *contaminated* — test items leak into training data [8, 17]. Second, they score a single answer and ignore what an agent would *do* with a belief. Third, they lack a strong, economically meaningful human baseline.

Real-world forecasting of *future* events dissolves all three problems at once. A scheduled sporting tournament provides a stream of questions whose answers (i) do not exist at query time, so contamination is impossible; (ii) arrive on a known date with unambiguous ground truth; (iii) are continuously priced by a liquid betting market that aggregates human and algorithmic beliefs into calibrated probabilities [28, 31]. The 2026 FIFA World Cup — 104 matches over five weeks, every one after current models’ training cutoffs — is close to an ideal such stream.

We use it to build **WC2026-Agents**, a benchmark that treats four frontier LLMs as forecasting agents and the betting market as a fifth competitor. For each match, each agent runs a fixed *search-act-reflect* loop: it searches the web for team news and prices, acts by returning calibrated 1X2 (team-A win / draw / team-B win) probabilities together with a virtual \$100 staking decision, and — once the score is known — reflects on its own forecast. We settle every bet at real 1X2 odds, so the benchmark scores three distinct capabilities that usually travel together but here come apart: *predictive calibration* (are the probabilities right?), *decision quality* (does the agent profit against a real price?), and *self-knowledge* (does it know when it was wrong?).

A reference evaluation of the four 2026 models illustrates why these axes matter. On *what will happen*, the agents are nearly interchangeable: they choose the same most-likely outcome in 92% of matches, their accuracies fall in a three-point band, and none beats the market’s Brier score — an efficient-market result that a single-answer benchmark would report as a null. On *what to do about it*, they diverge sharply: betting ROI spans -18% to +10%; fading the market loses money for every agent; the fraction of forecasts that even mention the market spans 12% to 100%; and, on their

own wrong picks, self-reported error rates span 36% to 86%. These behavioural axes, not accuracy, are where frontier models differ.

Contributions.

- **A contamination-free agentic-forecasting benchmark** over all 104 matches of a future World Cup: 416 forecasts and 414 reflections from four frontier LLMs, with verbatim reasoning, paired pre-/post-match structure, and a reproducible evaluation suite (§4, §5).
- **A market baseline as a fifth competitor.** We release per-match 1X2 odds and the settlement machinery so any forecaster can be scored not only for calibration but for realized betting profit against a real price (§4.5).
- **A curated, documented dataset.** Fixtures, team orientation, scores (including penalty shootouts and who advanced) were hand-resolved from four heterogeneous transcripts and cross-checked four ways (§4.4).
- **A reference evaluation** spanning calibration, betting/decision quality, reasoning content, and reflection honesty, with the findings above and a distillation of the agents' recurring blind spots into candidate features (§6).

2 Related Work

LLMs as forecasters. Retrieval-augmented LLMs can approach human crowd accuracy on judgemental questions [11], and LLM ensembles rival human crowds [23], echoing the human superforecaster literature [26]. Dynamic forecasting benchmarks such as FORECASTBENCH [15] and early neural event forecasting [37] score probabilistic accuracy. We add three elements those lack: a fixed agentic loop, a betting-based *decision* score against a real market, and paired self-reflection.

LLM agents, reasoning, and reflection. Agentic scaffolds interleave reasoning and tool use [30, 36] and let models critique and revise their own outputs [18, 24], though unaided self-correction is fragile [13]. Surveys chart the space [29, 34]. Our reflection prompt is a minimal, standardized probe of self-knowledge at tournament scale.

Calibration and uncertainty. Modern networks are often miscalibrated [10]; LLM confidence can be elicited but is imperfectly calibrated [14, 27, 35]. We measure calibration with reliability curves and expected calibration error [20], and full-distribution quality with the Brier score [4] and log-loss.

Betting and prediction markets. Bookmaker odds are strong, hard-to-beat probability estimates [1, 28, 32], consistent with market efficiency [7] and exhibiting known regularities such as the favorite-longshot bias [16, 25]. Machine-learning football models are a long line of work [3, 5]. We use the market both as a baseline and as the price an agent may follow or fade.

Benchmarks and contamination. Static benchmarks [12, 17] face contamination as models train on ever more of the web [8]. Forecasting strictly future events is contamination-free by construction, a property we exploit. Finally, algorithm aversion [6] and overconfidence [19] frame our behavioural analysis.

3 Task Formulation

A match x has an outcome $o(x) \in \mathcal{Y} = \{\text{team_a_win}, \text{draw}, \text{team_b_win}\}$, where team A is the first-listed side. A forecasting agent g observes x and its public information environment and returns a distribution $p_g(x) \in \Delta(\mathcal{Y})$ and a decision $d_g(x) = (\text{pick}, \text{stake})$ with stake $\in [0, 100]$. After the match it returns a reflection $r_g(x)$. The market m supplies an implied distribution $p_m(x)$ obtained by removing the vigorish from published 1X2 odds; it makes no bets and files no reflection, serving as a fifth forecaster and as the price at which agent bets settle. Knockout ties that finish level go to a shootout; because a 1X2 bet settles on the 90-minute result, we take $o(x)$ to be that result (a draw for such ties) and separately record the shootout winner (§4.4).

4 The WC2026-Agents Benchmark

4.1 Agent protocol

Each match generates two turns in a single per-model conversation. The *pre-match* turn ($T-24\text{h}$) instructs the agent that it is in a research study (virtual money; do not refuse), names the fixture, tells it to use web search, and requires a strict JSON reply:

```
{
  "probabilities": {"team_a_win", "draw",
    "team_b_win"},
  "bet": {"pick", "stake_usd"},
  "reasoning": ..., "key_sources": [urls]}
```

The *post-match* turn provides only the final score and asks — without web access — for a reflection: an outcome_vs_prediction label (correct/partially_correct/incorrect), a calibration self-judgement, a luck attribution, one key_factor_i_missed, one key_factor_i_overweighted, a counterfactual, a would-bet-differently verdict, and a confidence. The prompt text is byte-identical across agents; only the model changes.

4.2 Models

We evaluate four frontier assistants queried through their consumer interfaces with web tools enabled and held fixed across the tournament: Claude Opus 4.8 [2], ChatGPT (GPT-5.5 “Thinking”, high reasoning effort) [21], Gemini 3.1 Pro [9], and Grok in Expert Mode [33]. Using production assistants (rather than raw API base models) is deliberate: it captures the retrieval-and-reason *agent* a user actually deploys, which is the object of study.

4.3 Data collection

Between 11 June and 19 July 2026, each agent was run on all 104 fixtures ($m01-m72$ group, $m73-m104$ knockout), producing a pre-match forecast $\sim 24\text{h}$ before kickoff and a post-match reflection after the result. Raw transcripts are archived read-only. The pre-match window strictly precedes each outcome, so no target could appear in any model’s training data.

4.4 Curation and quality control

Group-stage transcripts follow a regular structure and were parsed with a string-aware JSON extractor. The knockout transcripts are not regex-parseable: they use four different header renderings,

Table 1: Dataset coverage. Two Gemini group reflections are absent in the source and flagged, not imputed.

Stage	Matches	Agents	Forecasts	Reflections	Bets placed
Group	72	4	288	286	232
Knockout	32	4	128	128	103
Total	104	4	416	414	335

some Round-of-16 fixtures were pasted with bracket placeholders (“W74 vs W77”) whose numbering does not match reality, and final scores appear in a dozen phrasings (penalty shootouts, lower-cased names, occasionally reversed team order). We therefore *hand-resolved* fixtures, team orientation, and ground truth for all 32 knockout matches from the transcripts, encoding them in a verified table. As an automatic guard, the human-pasted score line in each of the four agents’ transcripts is compared to that table; the four-way check flagged zero disagreements across all 128 knockout cells. Ground truth records the 90-minute result, the shootout result where applicable, and which side advanced. Every deviation from the JSON schema (recovered braces, duplicated pastes, two missing Gemini reflections) is logged rather than dropped.

4.5 Market odds

We release pre-match 1X2 odds for all 104 matches with a source URL each, oriented to team A and vig-removed to implied probabilities (mean overround 1.05). Knockout odds are opening lines from public sportsbook previews; we validated them against the market probabilities the agents themselves cited in their reasoning (e.g. our de-vigged Brazil–Japan line, 55/25/20, matches a cited “56/25/19”), and corrected one line whose implied hold was economically impossible. The market is thus both an independent baseline and the settlement price for agent bets.

4.6 Schema, statistics, accessibility

The release is one row per (agent, match): fixture and stage; ground truth with penalty and advancement fields; the pre-match probabilities, pick, stake, verbatim reasoning and cited URLs; the eight reflection fields; and parse-issue flags. A tidy 104-match table joins all agents and the market for analysis. Table 1 summarizes coverage. All data (CC BY 4.0) and code (MIT), including the one-command pipeline that regenerates every table and figure, are public; the dataset requires no access request.

An illustrative record. For match *m*92 (Brazil vs. Norway, Round of 16), one agent returned the pre-match forecast {team_a_win: 0.55, draw: 0.25, team_b_win: 0.20}, a \$40 stake on Brazil, and the reasoning “Brazil enters as heavy favorites with superior squad depth . . . Norway’s Haaland–Ødegaard threat and historical edge make an upset possible.” Norway won 2–1. The paired reflection then labelled the outcome incorrect, named the missed factor as “Norway’s historical dominance and counter-attacking effectiveness . . . amplified by Haaland,” and the over-weighted factor as “Brazil’s overall squad depth and pedigree.” Every record has this pre/post structure, letting a user trace a belief, a decision, a result, and a self-assessment for the same event.

Table 2: Task T1 — calibrated prediction over all 104 matches. Best in bold; $\uparrow/\downarrow$ = higher/lower better. The market attains the best Brier score.

	Acc. $\uparrow$	Brier $\downarrow$	LogLoss $\downarrow$	ECE $\downarrow$	Draw mass
Claude	.664	.4705	.806	.116	.234
ChatGPT	.683	.4729	.814	.111	.235
Gemini	.654	.4828	.820	.068	.238
Grok	.683	.4706	.803	.097	.227
Market	.683	.4688	.807	—	.229

Table 3: Task T2 — betting/decision quality over all 104 matches. Each agent settles its own picks at real 1X2 odds. “Contr.” = contrarian (against-market) bets.

	#Bets	Staked	Net \$	ROI %	Hit	Contr. share/ROI
Claude	73	1519	−275	−18.1	.329	.58 / −24%
ChatGPT	55	1471	+118	+8.0	.509	.36 / −14%
Gemini	103	8660	+322	+3.7	.631	.14 / −12%
Grok	104	6305	+650	+10.3	.673	.05 / +33%

5 Benchmark Tasks and Metrics

The suite scores three capabilities.

(T1) Calibrated prediction. Accuracy ($\arg \max p_g = o$); multi-class Brier $\sum_i (p_{g,i} - \mathbb{1}[o=i])^2$; log-loss $-\log p_{g,o}$; and, for the top pick, a reliability curve and expected calibration error (ECE) [20] over ten confidence bins. *Convergence* — the fraction of matches on which all agents share an $\arg \max$ — quantifies how much accuracy can separate them.

(T2) Decision quality. Each agent’s own pick/stake settles at the match’s decimal odds on the 90-minute result: profit = stake $\times$ (dec − 1) on a win, else −stake. We report bankroll, ROI, hit-rate, and a *contrarian* split: a bet is contrarian if its pick differs from the market’s most-likely outcome, isolating the value of fading the price. The market’s implied probabilities are scored under T1 as a fifth forecaster.

(T3) Behavioural transparency. Every reasoning and reflection field is coded into 13 factor categories with a published keyword lexicon (no model in the loop), giving each agent’s information diet and its *self-knowledge*: on the matches its $\arg \max$ got wrong, how often does its reflection label itself incorrect?

6 Reference Evaluation

We evaluate the four 2026 agents and the market. Tables 2 and 3 give the headline numbers; figures expand each result.

6.1 The agents converge; accuracy cannot separate them

The four agents issue an identical top pick in 92% of matches (91.7% group, 93.8% knockout; Fig. 1). Pairwise agreement reaches 100% (ChatGPT–Grok) and never falls below 92%. Accuracy therefore spans only three points (Table 2), i.e. three matches. Figure 2 shows the mechanism: each agent’s team-A probability tracks the market’s

almost perfectly ($r = 0.97\text{--}0.99$), so the agents are, to first order, reading and lightly re-expressing the same price. On the eight matches where they are not unanimous, the split is always 3–1 (a single agent dissents), and the dissenter is usually Gemini (six of eight) — consistent with it being the least market-anchored agent (Fig. 6). All eight are near-coin-flips where the market itself is close to even, so a shade more weight on one side tips the arg max. Disagreement thus appears exactly where the signal is weakest, not where an agent has found private information.

6.2 No agent beats the market

Scored as a forecaster, the market attains the best Brier score of all five competitors (0.469 vs. 0.471 for the best agent) and ties the best accuracy (Table 2; per-stage in Table 4). It also identifies the advancing side in 24 of 32 knockout ties (75%), equal to the agents. This is the expected efficient-market outcome [28]: web-enabled agents recover the market’s information but do not systematically exceed it. The tournament was mildly favorite-friendly, which the released odds let users control for: binning matches by the market’s implied probability for its favored side, that side won 62% of the time when priced near 48% and 78% when priced near 62%, while heavy favorites (~81%) were well-calibrated. In a cup with slightly fewer upsets than priced, every forecaster that respected the market — agents and market alike — looks under-confident in hindsight (§6.3), a difficulty profile the benchmark exposes rather than hides.

6.3 Calibration does separate them

All four agents lie *above* the reliability diagonal: they are **under-confident** on their top pick, which wins more often than its stated probability (Fig. 3). Gemini is best calibrated (ECE 0.068), then Grok (0.097), ChatGPT (0.111), Claude (0.116). Calibration and full distribution can disagree: Gemini has the best top-pick ECE but the worst Brier, because it spreads probability mass more aggressively and stakes larger (Table 3).

6.4 Decision quality diverges where accuracy does not

Settling each agent’s own picks at real odds, betting ROI spans –18.1% (Claude) to +10.3% (Grok), with net results of –\$275 to +\$650 (Table 3, Fig. 4). The spread is not driven by who they pick — they pick alike — but by *how much they stake and whether they respect the price*. Gemini stakes most aggressively (\$8,660 across 103 bets); Claude stakes least and most often on draws and underdogs. As a reference, a naive baseline that stakes a flat \$100 on the market favourite every match returns +\$1,041 (grey line, Fig. 4), out-earning all four agents in absolute terms at comparable ROI — a betting counterpart of the result that no agent beats the market. (Favorites over-performed their price this cup, §6.3; the released odds let users hold this difficulty fixed.)

6.5 Fading the market is unprofitable for every agent

Decomposing bets by their relationship to the price explains the P&L (Fig. 5). Contrarian (against-market) bets hit 21–40% and return –24% to +33% (the lone positive being Grok’s five such bets), while

market-conforming bets are the source of all profit. The ordering is monotone: the more an agent fades the market, the worse it does. Claude cites the market in every forecast yet bets against it 58% of the time — it *sees* the price and overrides it — and is the only agent with a net loss.

6.6 Staking discipline is a distinct axis

Beyond *which* side to back, agents differ in how they *size* a bet. We correlate each agent’s stake with the edge it believed it held (its probability minus the market’s for the picked side) and with its confidence. Grok’s stake tracks its confidence most tightly (Pearson $r=0.85$) and its claimed edge positively ($r=0.13$) — Kelly-like behaviour that concentrates money on its strongest reads. Gemini and ChatGPT also stake up with confidence ($r=0.54$ and 0.53), but Gemini’s mean stake (\$84 of a \$100 budget) is a near-max-bet habit that inflates variance — its bankroll in Fig. 4 is the most volatile of the four. Claude is the reverse: its stake is *negatively* related to confidence ($r = -0.20$), because it sizes up precisely on the low-confidence, against-market “value” plays that tend to lose. Staking discipline is therefore a separate, measurable axis, only loosely coupled to whether a pick is correct, and the released stakes let others re-score the agents under any fixed staking policy.

6.7 Information diets differ while picks converge

Coding the reasoning reveals sharply different information diets that nonetheless yield the same picks (Fig. 6). The share of forecasts that cite the betting market or odds ranges from **100% (Claude) and 92% (ChatGPT) to 12% (Gemini)**; Gemini reasons almost entirely from football (attacking threat, defensive organisation, form) while barely referencing the price. This is the behavioural counterpart of Fig. 2: Gemini is market-blind in its *stated reasoning* yet still lands near the market, and consequently fades it in only 14% of its *bets*. Grok most often invokes “paper strength” (squad depth 88%, form 76%, FIFA ranking 61%); Claude most often invokes knockout variance (53%).

6.8 Reflection is a behavioural fingerprint

Given only the final score, the agents reflect very differently on their own mistakes (Fig. 8). On the picks each got wrong, the rate of admitting “incorrect” ranges from **86% (Gemini) to 36% (ChatGPT)**, which softens 55% of its errors to “partially correct”; Claude (49%) and Grok (61%) sit between. When they were right, the rate of claiming “correct” ranges from 68% (Claude) to 94% (Grok). Reflection style is thus a stable per-agent trait, orthogonal to accuracy — relevant wherever a model’s self-report gates a downstream action [13, 24].

6.9 Shared blind spots and candidate features

Eight matches were mispredicted by *all four* agents (Fig. 7); in every case they backed the favourite. Four were penalty ties in which a heavy favourite (Germany, 75% consensus) failed to win in 90 minutes; the rest were regulation upsets (e.g. Norway over Brazil, Belgium 4–1 over the USA, Spain over France). The agents’ own post-mortems are strikingly consistent: across all wrong picks, the

Agreement of the four agents' top pick




Figure 1: Convergence. Left: pairwise same-pick rate. Right: match counts by consensus; genuine disagreement is rare and unanimous picks are right about twice as often as wrong.

Agent vs. market-implied probability (team-A win), 104 matches per agent




Figure 2: Each agent's probability for a team-A win against the market-implied probability, all 104 matches. Points hug the diagonal (r shown); Gemini is the loosest, consistent with it rarely citing the market (Fig. 6).

factor most often named as *missed* is underdog defensive organisation (62%), then knockout/penalty variance (32%); the factor most often named as *over-weighted* is the favourite's squad quality (40%) and attacking firepower (38%). These yield concrete, minable priors for future systems: down-weight a favourite's paper talent against a disciplined low block, and explicitly price the draw→shootout path in knockout ties.

7 Discussion and Impact

The benchmark separates three capabilities that a single-answer score conflates. On prediction the 2026 agents are interchangeable and no better than the market; on *decision* and *self-knowledge* they differ by large, stable margins. For practitioners the lesson is that selecting an LLM for a consequential forecasting-and-acting

pipeline should weigh calibration, staking discipline, and honest self-assessment, not headline accuracy. Because the protocol is a thin wrapper around a consumer assistant, WC2026-Agents is a template: the same search–act–reflect loop and market settlement apply to any scheduled, market-priced event stream (other tournaments, elections, product launches), enabling live, rolling, contamination-free evaluation and longitudinal tracking of model releases. We release the loop, lexicon, and settlement code to that end.

Two results deserve emphasis for benchmark design. First, the convergence of four independently built models cautions against reading a single leaderboard number as a capability signal: when models share a retrieval surface they also share errors, so an ensemble inherits rather than cancels the common bias — a limit on

Reliability of the top pick — colour = agent, grey = market (all under-confident)




Figure 3: Reliability of each agent’s top pick, with the market as a grey reference curve. All lie in the shaded (under-confident) region; marker area $\propto$ bin count. ECE also appears in Table 2.

the “wisdom of the silicon crowd” [23] when the crowd reads the same web. Second, that no agent beats a vig-bearing market is the sharpest available null for an over-eager “LLMs can predict the future” narrative: the market is not a weak baseline but a near-efficient [7, 28] aggregator that the agents match but do not exceed, so a benchmark that omits it would mistake market-level performance for super-human forecasting. Surfacing such failure modes is precisely the role of a datasets-and-benchmarks contribution.

8 Using WC2026-Agents

The release supports several lines of work out of the box, each a single join on the tidy 104-match table. **(i) Calibration and uncertainty.** 416 forecasts with confidences and outcomes, plus a market reference, form a contamination-free testbed for LLM (mis)calibration and confidence elicitation [27, 35]. **(ii) Agentic decision-making.** Paired probabilities, stakes, and real odds let researchers score arbitrary staking policies (flat, Kelly, learned) and quantify the value of following versus fading a price — a decision layer absent from prediction-only benchmarks [15]. **(iii) Self-knowledge.** 414 paired reflections let one test whether models know when they were wrong [13, 14] and whether self-reports predict error, using our outcome_vs_prediction labels against realized correctness. **(iv) Reasoning and evidence use.** Verbatim reasoning with cited URLs supports factor-attribution and source-use studies via the released lexicon. **(v) Temporal robustness and longitudinal tracking.** Because every target post-dates training, the data probes genuine forecasting rather than recall; the thin

search-act-reflect protocol re-runs on any future tournament, enabling live, rolling evaluation of successive model releases against a moving, uncontaminated target.

9 Ethics

All betting in this study is virtual; the dataset is a measurement instrument, not gambling advice, and we report profit only to score decision quality. No human subjects or personal data are involved: the records are model outputs and public fixtures, scores, and odds. Odds come from public preview articles with source URLs; we take no book’s proprietary feed. A foreseeable misuse — treating any agent’s edge as an actionable betting strategy — is unsupported by our data, which show no agent beats the market. We disclose that generative AI assisted the parsing and drafting of this work; all findings were verified against the released data by the authors.

10 Limitations

This is a single tournament (104 matches); per-agent differences have wide intervals and we stress patterns over precise rankings. Agents were queried through consumer interfaces with tools enabled, so retrieval and decoding are not fully controlled and providers may update models mid-window — we hold the user-visible model fixed and record it. Odds are opening lines from mostly one book, not closing consensus. Reflections are the model’s own free text, coded by a transparent but lexical scheme. These are the natural costs of measuring deployed agents on real, un-repeatable events, and the released pipeline makes every choice auditable.

11 Conclusion

WC2026-Agents turns a future World Cup into a contamination-free benchmark of LLMs as forecasting agents, paired with the betting market as a fifth, economically grounded competitor. Its reference evaluation shows that frontier agents predict alike and cannot beat the market, yet diverge sharply in calibration, betting discipline, and the honesty of their self-reflection — the axes on which they actually differ. We release the data, odds, and evaluation suite to support live, rolling evaluation of agentic forecasting.




Figure 4: Cumulative virtual profit over the 104-match tournament, each agent settling its own stakes at real 1X2 odds. Trajectories diverge despite nearly identical point predictions.

Fading the market loses for every agent




Figure 5: Against-the-market analysis. Top: share of each agent's bets that faded the market favourite. Bottom: net profit from market-conforming vs. contrarian bets. Fading the market loses money for all four agents.




Figure 6: Share of pre-match forecasts citing each factor, per agent (dumbbell: one marker per agent). Market/odds citation ranges from 12% to 100%.




Figure 7: The eight matches mispredicted by all four agents. Every agent backed the favourite; four ties were lost on penalties after a 90-minute draw.




Figure 8: Self-assessment on each agent's own wrong picks. Owning error ("incorrect") ranges from 36% to 86% across agents.

References

- [1] Giovanni Angelini and Luca De Angelis. 2019. Efficiency of online football betting markets. *International Journal of Forecasting* 35, 2 (2019), 712–721.
- [2] Anthropic. 2026. Claude Opus 4.8. <https://www.anthropic.com/claude>.
- [3] Rahul Baboota and Harleen Kaur. 2019. Predictive analysis and modelling football results using machine learning approach for English Premier League. *International Journal of Forecasting* 35, 2 (2019), 741–755.
- [4] Glenn W. Brier. 1950. Verification of forecasts expressed in terms of probability. *Monthly Weather Review* 78, 1 (1950), 1–3.
- [5] Anthony C. Constantinou, Norman E. Fenton, and Martin Neil. 2012. pi-football: A Bayesian network model for forecasting Association Football match outcomes. *Knowledge-Based Systems* 36 (2012), 322–339.
- [6] Berkeley J. Dietvorst, Joseph P. Simmons, and Cade Massey. 2015. Algorithm aversion: People erroneously avoid algorithms after seeing them err. *Journal of Experimental Psychology: General* 144, 1 (2015), 114–126.
- [7] Eugene F. Fama. 1970. Efficient capital markets: A review of theory and empirical work. *The Journal of Finance* 25, 2 (1970), 383–417.
- [8] Shahriar Golchin and Mihai Surdeanu. 2024. Time travel in LLMs: Tracing data contamination in large language models. In *International Conference on Learning Representations (ICLR)*.
- [9] Google DeepMind. 2026. Gemini 3.1 Pro model card. <https://deepmind.google/models/model-cards/gemini-3-1-pro/>.
- [10] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. On calibration of modern neural networks. In *International Conference on Machine Learning (ICML)*, 1321–1330.
- [11] Danny Halawi, Fred Zhang, Chen Yueh-Han, and Jacob Steinhardt. 2024. Approaching human-level forecasting with language models. *arXiv preprint arXiv:2402.18563* (2024).
- [12] Dan Hendrycks, Collin Burns, Steven Basart, et al. 2021. Measuring massive multitask language understanding. In *International Conference on Learning Representations (ICLR)*.
- [13] Jie Huang, Xinyun Chen, Swaroop Mishra, et al. 2024. Large language models cannot self-correct reasoning yet. In *International Conference on Learning Representations (ICLR)*.
- [14] Saurav Kadavath, Tom Conerly, Amanda Askell, et al. 2022. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221* (2022).
- [15] Ezra Karger, Houtan Bastani, Chen Yueh-Han, et al. 2024. ForecastBench: A dynamic benchmark of AI forecasting capabilities. *arXiv preprint arXiv:2409.19839* (2024).
- [16] Steven D. Levitt. 2004. Why are gambling markets organised so differently from financial markets? *The Economic Journal* 114, 495 (2004), 223–246.
- [17] Percy Liang, Rishi Bommasani, Tony Lee, et al. 2023. Holistic evaluation of language models. *Transactions on Machine Learning Research (TMLR)* (2023).
- [18] Aman Madaan, Niket Tandon, Prakhar Gupta, et al. 2023. Self-Refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [19] Don A. Moore and Paul J. Healy. 2008. The trouble with overconfidence. *Psychological Review* 115, 2 (2008), 502–517.
- [20] Mahdi Pakdaman Naeini, Gregory F. Cooper, and Milos Hauskrecht. 2015. Obtaining well calibrated probabilities using Bayesian binning. In *AAAI Conference on Artificial Intelligence*.
- [21] OpenAI. 2026. GPT-5.5. <https://openai.com/index/gpt-5-5-instant/>.
- [22] Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, et al. 2023. Generative agents: Interactive simulacra of human behavior. In *ACM Symposium on User Interface Software and Technology (UIST)*.
- [23] Philipp Schoenegger, Indre Tuminauskaitė, Peter S. Park, and Philip E. Tetlock. 2024. Wisdom of the silicon crowd: LLM ensemble prediction capabilities rival human crowd accuracy. *Science Advances* 10, 45 (2024).
- [24] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [25] Erik Snowberg and Justin Wolfers. 2010. Explaining the favorite–long shot bias: Is it risk-love or misperceptions? *Journal of Political Economy* 118, 4 (2010), 723–746.
- [26] Philip E. Tetlock and Dan Gardner. 2015. *Superforecasting: The Art and Science of Prediction*. Crown.
- [27] Katherine Tian, Eric Mitchell, Allan Zhou, et al. 2023. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models. In *Empirical Methods in Natural Language Processing (EMNLP)*.
- [28] Erik Strumbelj. 2014. On determining probability forecasts from betting odds. *International Journal of Forecasting* 30, 4 (2014), 934–943.
- [29] Lei Wang, Chen Ma, Xueyang Feng, et al. 2024. A survey on large language model based autonomous agents. *Frontiers of Computer Science* 18, 6 (2024).
- [30] Jason Wei, Xuezhi Wang, Dale Schuurmans, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [31] Justin Wolfers and Eric Zitzewitz. 2004. Prediction markets. *Journal of Economic Perspectives* 18, 2 (2004), 107–126.
- [32] Fabian Wunderlich and Daniel Memmert. 2018. The betting odds rating system: Using soccer forecasts to forecast soccer. *PLOS ONE* 13, 6 (2018), e0198668.
- [33] xAI. 2026. Grok Expert Mode. <https://x.ai>.
- [34] Zhiheng Xi, Wenxiang Chen, Xin Guo, et al. 2023. The rise and potential of large language model based agents: A survey. *arXiv preprint arXiv:2309.07864* (2023).
- [35] Miao Xiong, Zhiyuan Hu, Xinyang Lu, et al. 2024. Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation. In *International Conference on Learning Representations (ICLR)*.
- [36] Shunyu Yao, Jeffrey Zhao, Dian Yu, et al. 2023. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*.
- [37] Andy Zou, Tristan Xiao, Ryan Jia, et al. 2022. Forecasting future world events with neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*.

A Reproducibility

One command (`python src/run_all.py`) rebuilds every processed table, analysis CSV, and figure from the raw transcripts and odds; the paper’s numbers are read directly from those outputs. The repository documents the schema, the 13-category reasoning lexicon, per-match ground truth with penalty and advancement fields, and all parse-issue flags. Randomness is not used in evaluation. The four-way score cross-check and the odds/citation validation are part of the build and re-run on every invocation.

B Per-stage metrics

Table 4 breaks T1 down by phase. Knockout matches are markedly more predictable for every competitor (accuracy .719 vs. $\sim$.65; lower Brier), because seeding concentrates mismatches; the four agents are again indistinguishable within each phase, and the market again leads on Brier. The full 13×4 factor table, the reliability bins behind Fig. 3, the contrarian breakdown behind Fig. 5, and the match-level upset details are released as CSVs under `data/analysis/` and regenerated by `src/run_all.py`.

Table 4: Task T1 by phase (accuracy / Brier / log-loss). Group ($n=72$) and knockout ($n=32$).

	Group			Knockout		
	Acc	Brier	LogL	Acc	Brier	LogL
Claude	.639	.489	.828	.719	.428	.756
ChatGPT	.667	.493	.838	.719	.429	.759
Gemini	.625	.485	.820	.719	.477	.822
Grok	.667	.491	.826	.719	.425	.752
Market	.667	.486	.827	.719	.431	.761